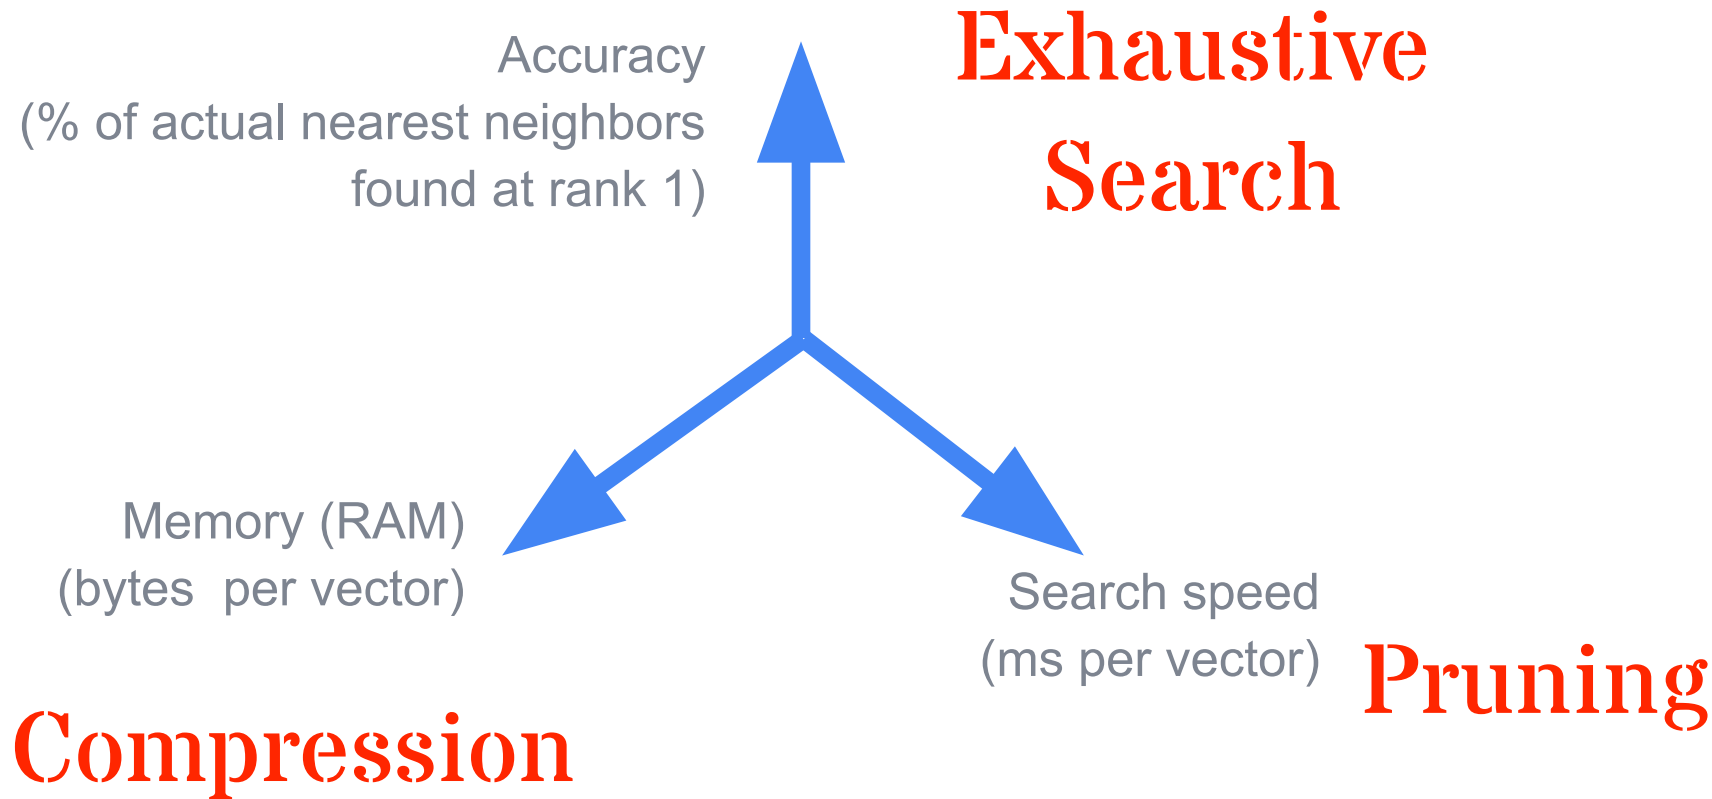


Vector search #8 – Quantization for lossy vector compression

Tradeoffs of vector search

3-handed tradeoff



In this class

- Mainly about the “compression” hand
 - Fix size of representation
 - Because RAM is constrained
 - Operating points between the two other hands
-
- Examples from Faiss
 - Implements many index types
 - Explore boundaries

Next page: Faiss wiki


<https://github.com/facebookresearch/faiss/wiki/Indexing-1G-vectors>

Memory budget: max
32 bytes per vector

Speed:
Queries per
second
(Log scale)

Each curve
= one type
of index

Accuracy:
1-recall@1



Built in conjunction with a notebook

- This sign:



- Means there is related content in the notebook.

Vector quantization

Vector quantization: definition

- Quantization: map a vector to an integer
 - Input is (supposed to be) continuous
 - Output is discrete
- Integer \equiv bit array of fixed size \equiv code
- Reconstruction: inverse map
 - We recover only an approximation
 - The reconstruction is lossy
- Evaluation: Mean Squared Error
 - Because it has nice arithmetic properties...
 - Invariant with d-dim rotation

$$q : \mathbb{R}^d \mapsto \{0, \dots, k - 1\}$$

$$r : \{0, \dots, k - 1\} \mapsto \mathbb{R}^d$$

$$\text{MSE} = \mathbb{E}_x [\|r(q(x)) - x\|^2]$$


Codes and relationship with clustering

- Numbering is sequential
 - Otherwise do a mapping in the discrete domain
- Size of codes $\lceil \log_2(k) \rceil$

- Quantization cell
 - Locus of vectors that produce the same code

$$q^{-1}(\{i\}), \quad i \in \{0, \dots, k-1\}$$

- All quantization cells are a partition of input space
- From a discrete point of view: clustering
 - Reconstruction values are called “centroids”



Lloyd's optimality conditions (reminder)



- To minimize the MSE for discrete sets
 - 1: a vector should be assigned to the nearest reconstruction
 - Otherwise re-assign that vector: decrease MSE!
 - 2: each centroid should be the center of mass of points assigned to it
 - Otherwise just move the center of mass: decrease MSE!
 - Necessary, not sufficient
-
- The k-means algorithm
 - Iterate the two optimality conditions

Cases where the optimal centroids are known


- Uniform distribution
 - A_x Lattice
- Uniform over a sphere
 - Integer lattice on sphere
- Unknown for Gaussian data

[Paulevé et al, Locality sensitive hashing: a comparison of hash function types and querying mechanisms, Pattern recog letters 2010]


[Sablayrolles et al, spreading vectors for similarity search, ICLR'19]



(b) A_2 lattice



(c) k-means
Uniform distribution



(d) k-means
Gaussian distribution

Figure 3: Voronoi regions associated with random projections (a), lattice A_2 (b) and a k-means quantizer (c,d).

How optimal is k-means?

Optimality of k-means: practical considerations

- On small scale, practical k-means is quite off from the global optimum
- On large scale we don't know!
 - NP-hard is very hard


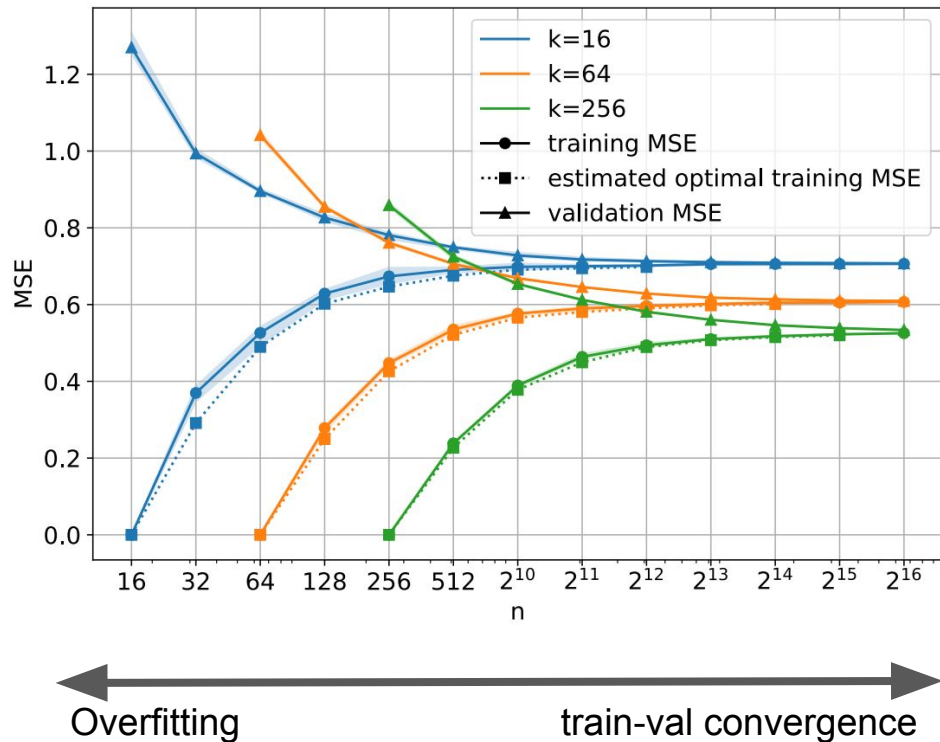


Figure 3: Histogram of MSE results for 1000 runs of k-means with $(n, k) = (24, 8)$. The experiment is run on 5 subsets of the same data distribution, hence the 5 plots. The exact global optimum is indicated in red. The estimate of this optimum from the k-means runs is the red dashed line.

[Augmented k-means, Touvron, Douze, Jégou, unpublished]




Difference between training and validation

- MSE can be computed on
 - Training vectors → Lloyd's conditions
 - Validation → how it's going to be used in reality
- Small and large-data regime
- Relevant parameter is n/k
 - Nb training points per centroid



Optimality of k-means: practical considerations


- As a function of the Points Per Centroid (ppc)
- Validation MSE
- Line thickness: min and max of 10 k-means runs



- It does not matter to have many initializations when k increases

Making good use of FLOPS

- Training a 100-centroid k-means on 1M vectors
 - Waste of resources?
- Given a certain computational budget
 - Balance nb of iterations and nb points per centroid



Scaling k-means

Scaling k-means

- Complexity
- Required to get larger codes
 - 3 bytes = 24 bits = 16M centroids
 - Not a very large code...
 - But a HUGE number of centroids
- Expensive stage is assignment
 - NNS problem
- Brute-force hardware scaling
 - CPU training
 - GPU training
 - Distributed training...



Matthijs Douze

March 8, 2019 · 📷

...

KMeans of 500M points to 10M centroids

144 dim, 20 iterations, 10x8 GPUs: 22h.

This is the largest k-means optimization we have done so far with Faiss. It required a version where the training set is distributed over 10 machines. The k-means logic was also re-implemented in Python. Nothing fancy, just brute force.

[Nistér Stewenius, Scalable recognition with a vocabulary tree, CVPR'06]
[Muja, Lowe, Scalable Nearest Neighbor Algorithms
for High Dimensional Data, PAMI'14]

Hierarchical k-means

- Inspired by bisection in 1D
- Run k-means recursively to subdivide
 - Iterate
- Often used for search
 - Basis of the FLANN library
- But sub-optimal in terms of MSE
 - See notebook

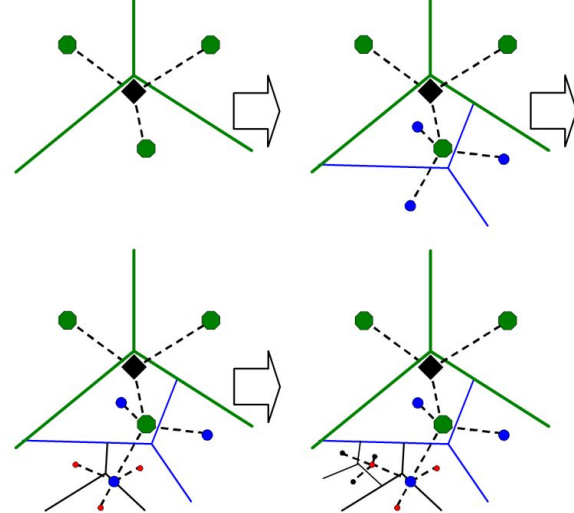


Figure 2. An illustration of the process of building the vocabulary tree. The hierarchical quantization is defined at each level by k centers (in this case $k = 3$) and their Voronoi regions.





Fig. 3. Projections of priority search k-means trees constructed using different branching factors: 4, 32, 128. The projections are constructed using the same technique as in [26], gray values indicating the ratio between the distances to the nearest and the second-nearest cluster centre at each tree level, so that the darkest values (ratio ≈ 1) fall near the boundaries between k-means regions.

Vector quantization for search: The inverted file

The Inverted File



- Cluster the space into k clusters of vectors
 - assign vectors to nearest centroid
 - Aka. “coarse quantizer”
- index = inverted list structure
 - maps centroid id \rightarrow list of vectors assigned to it
- search procedure:
 - 1. find $np \ll k$ nearest centroids to query vector
 - 2. scan the lists corresponding to the np centroids
- Objective: reduce the number of distance computations!





The Inverted File: number of centroids

- Tradeoffs

- For high-accuracy regime it is not necessary to be very selective: small nb clusters
- For low-accuracy regime it is better to filter more with clusters


- As a function of database size:

- Coarse quantization: k distance computations
- Scanning: (assuming balanced clusters): $n_{\text{probe}} * n / k$
- Exercise: find optimal k as a function of n for fixed n_{probe}



Comparing coarse quantizers

- Useful with 2 levels
 - Reduces complexity



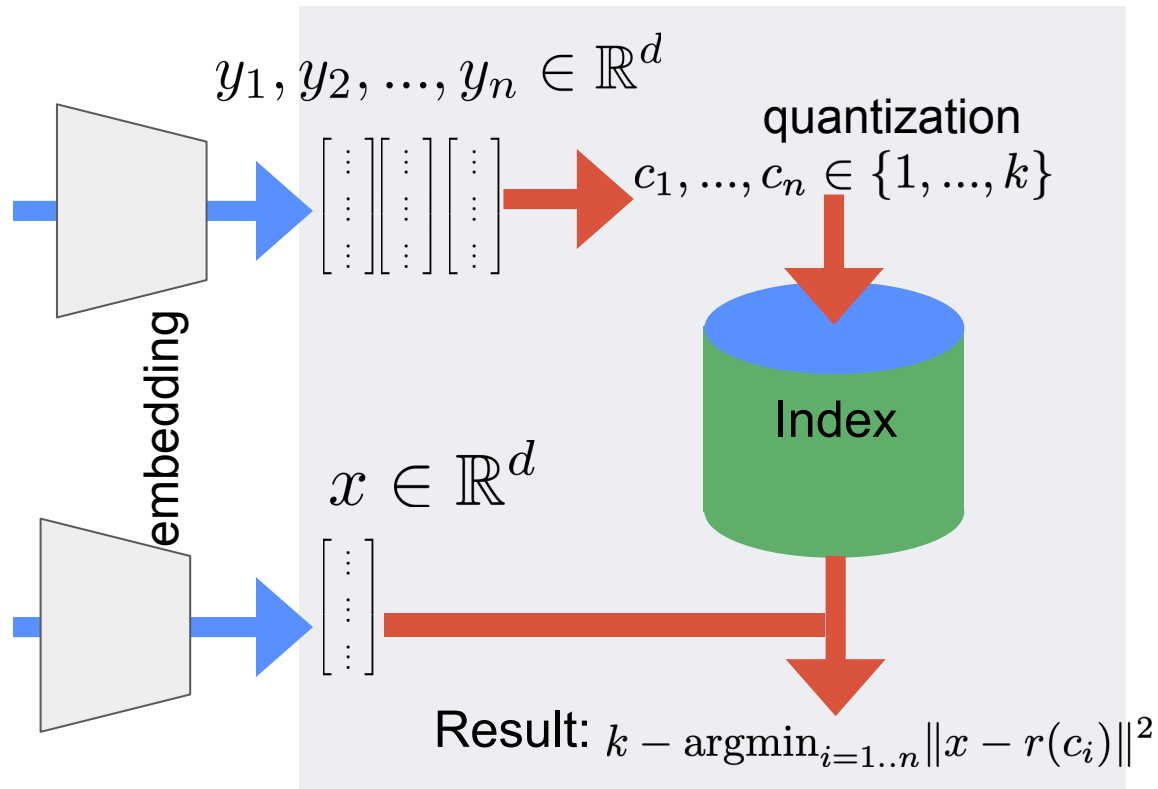
Vector quantization for compression

Compression and search: asymmetric case

Collection:




Query:




Symmetric vs. asymmetric comparison

- Consider asymmetric setting
 - no constraint on storage of query
 - keep full query vector, encode database vectors $q(y)$
- Distance estimator
 - reproduction value of the quantizer: centroid
 - approximate distance
$$\|x - y\| \approx \|x - q(y)\|$$



symmetric case



asymmetric case

- approximate nearest neighbor
$$\operatorname{argmin}_i \|x - q(y_i)\|$$



Distance computations with look-up tables

- For a given query x , there are k possible distances
 - Precompute a table!
 - At search time, look up the distances in the table.
 - No computations at search time, only look-ups
 - Useful if nb centroids \ll nb database elements



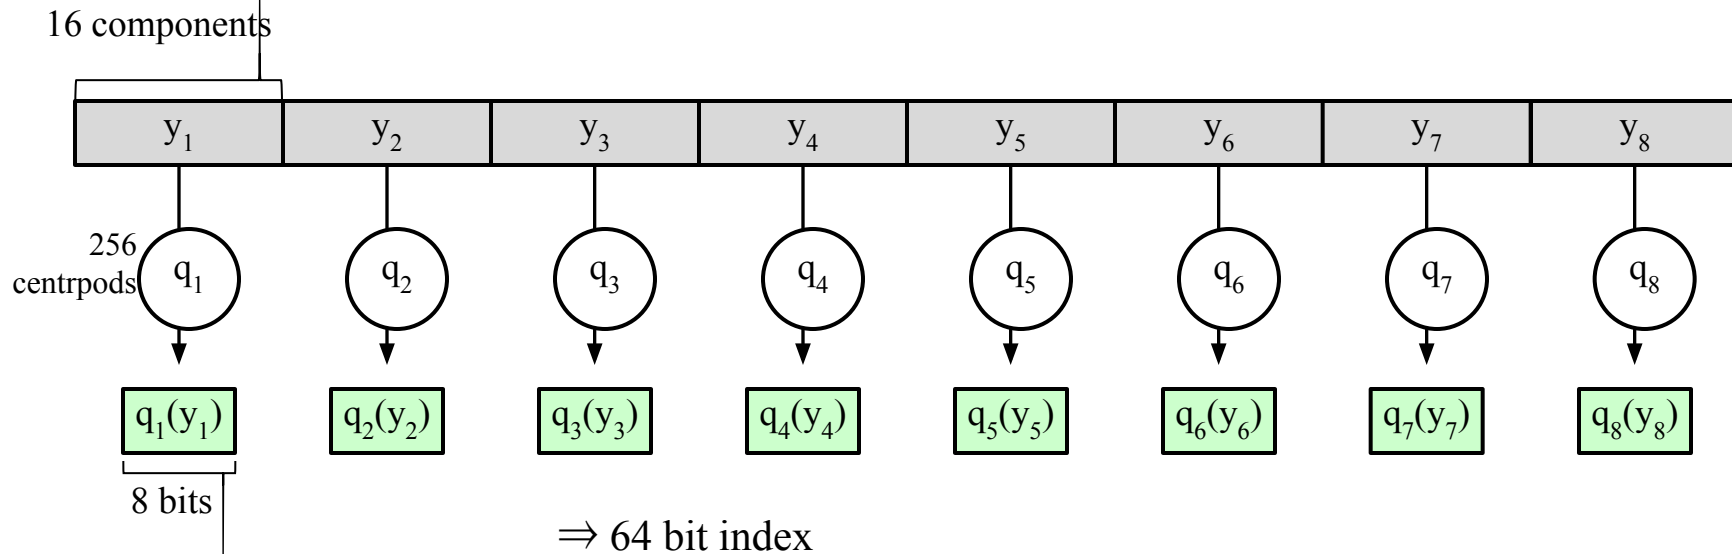
- However, this is pretty limited
 - Small codes – limited recall...

Multi-codebook quantization

Multiple-codebook quantization

- Combine multiple quantizers
 - Each has its codebook
 - Separate codes, total size $\sum_{m=1}^M \lceil \log_2(k_m) \rceil$
- In the following:
 - Product quantization
 - Additive quantization

Product Quantization



Multiple-codebook quantization

- reconstruction value: concatenation of centroids

$$y \approx [q_1(y^1), \dots, q_m(y^m)]$$

- distance computation: distance is additive

$$\|x - y\|^2 \approx \sum_{j=1}^m \|x^j - q_j(y^j)\|^2$$

- precompute M look-up tables!

Sizes & flops

	no compression	vector quantizer	product quantizer
code size	d	$\log_2(k)$	$m \cdot \log_2(k)$
quantization cost	N/A	$k \cdot d$	$k \cdot d$
distance computation cost	d multiply-adds	one look-up, one add	m look-ups, m adds
number of distinct values	N/A	k	k^m



Product Quantizer tradeoffs

- For a given code size
 - $\text{Code_size} = M * \log_2(k)$
- Higher k (and lower M)
 - Better accuracy
 - Larger quantization tables (slower)




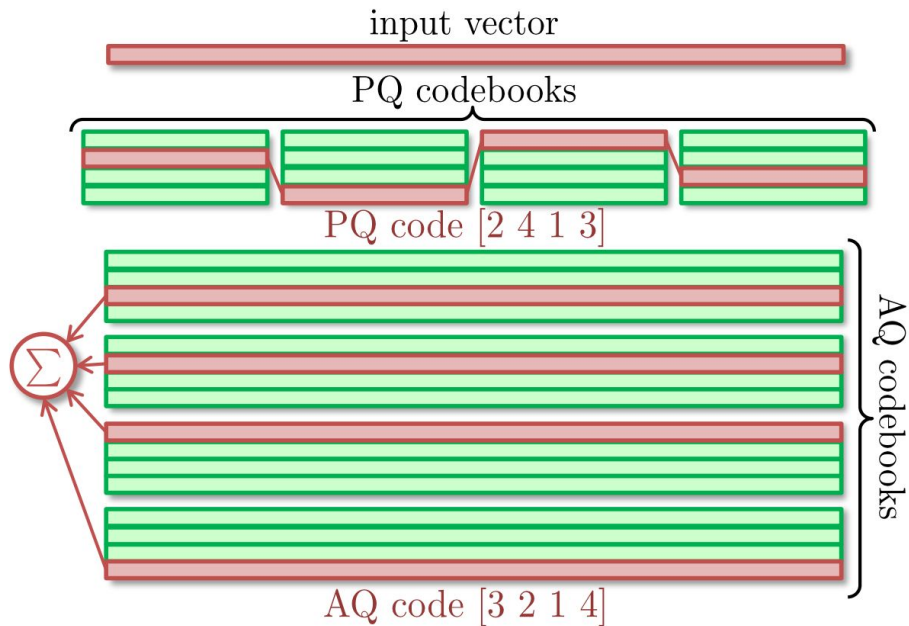
Fig. 1. SIFT: quantization error associated with the parameters m and k^* .

Optimized product quantizat

Additive quantization

Residual quantization

- Use multiple codebooks
 - But in dimension d



General additive quantization – LSQ

- 12x3 h
- + guest talks
- Evaluation via

Fast search with additive quantization

- 12x3 h
- + guest talks
- Evaluation via

Storing the norms

- 12x3 h
- + guest talks
- Evaluation via

Neural quantization methods

The catalyzer

- 12x3 h
- + guest talks
- Evaluation via

UNQ

- 12x3 h
- + guest talks
- Evaluation via

Polysemous codes

Various decoders

- 12x3 h
- + guest talks
- Evaluation via

Equivalence of decoders

- 12x3 h
- + guest talks
- Evaluation via

Practical implementation: IVFPQ

Combination of IVF and PQ

- PQ encodes residual

Look-up tables

- Predefine table
-

Parallelization strategies

- Parallelize over queries
- Parallelize intra query
- Coarse and fine quant

END